

ASSIGNMENT 1

Deadline: 11:59 pm, March 3, 2025

Submit via Blackboard with VeriGuide receipt.

Please follow the course policy and the school's academic honesty policy.

Submission Instructions:

Students must submit a zip file containing:

- A PDF file with their written solutions.
- The implemented `Assignment1_Q2.py` file.

For Problem 2, students must also either paste the output or include a screenshot of the output in the PDF file after executing their code.

1. Consider a feedforward neural network for text classification with 2-dimensional input vectors $\mathbf{x} = [x_1, x_2]$, a hidden layer of 3 ReLU neurons, and a 2-dimensional softmax output layer. The network uses cross-entropy loss and has no biases ($\mathbf{b}_1 = \mathbf{0}, \mathbf{b}_2 = \mathbf{0}$). Weight matrices are $\mathbf{W} \in \mathbb{R}^{3 \times 2}$ (input-to-hidden) and $\mathbf{V} \in \mathbb{R}^{2 \times 3}$ (hidden-to-output). Given a training example (\mathbf{x}, y) with $\mathbf{x} = [1, -1]$ and true label $y = 0$ (0-indexed). The ReLU activation function for the hidden layer is defined as:

$$\text{ReLU}(x) = \max(0, x).$$

The softmax function for the output layer is defined as:

$$\hat{y}_k = \frac{\exp(z_k)}{\sum_{i=0}^1 \exp(z_i)} \quad \text{for } k \in \{0, 1\} \quad \text{where } z = \mathbf{V}\mathbf{h}.$$

(1) Loss Computation (20 points)

Compute the hidden layer activations \mathbf{h} , output probabilities $\hat{\mathbf{y}}$, and cross-entropy loss L , given the following weights

$$\mathbf{W} = \begin{bmatrix} 0.4 & -0.1 \\ -1 & 0.6 \\ 0.5 & 1 \end{bmatrix}, \quad \mathbf{V} = \begin{bmatrix} 0.8 & 1.2 & -0.5 \\ 0.7 & -0.9 & 0.3 \end{bmatrix}.$$

(2) Gradient Backpropagation (20 points)

Derive the gradients of L with respect to all weights in \mathbf{W} and \mathbf{V} . Present results as:

$$\frac{\partial L}{\partial \mathbf{V}} = \begin{bmatrix} \frac{\partial L}{\partial V_{11}} & \frac{\partial L}{\partial V_{12}} & \frac{\partial L}{\partial V_{13}} \\ \frac{\partial L}{\partial V_{21}} & \frac{\partial L}{\partial V_{22}} & \frac{\partial L}{\partial V_{23}} \end{bmatrix}, \quad \frac{\partial L}{\partial \mathbf{W}} = \begin{bmatrix} \frac{\partial L}{\partial W_{11}} & \frac{\partial L}{\partial W_{12}} \\ \frac{\partial L}{\partial W_{21}} & \frac{\partial L}{\partial W_{22}} \\ \frac{\partial L}{\partial W_{31}} & \frac{\partial L}{\partial W_{32}} \end{bmatrix}.$$

Explicitly show the chain rule steps for backpropagation.

2. In this assignment, you will implement core components of neural networks: a simple feed-forward Multi-Layer Perceptron (MLP) and the Cross Entropy Loss function. This will help you understand the fundamental building blocks of deep learning models.

(1) Implementing MLP Forward Pass (30 points)

Complete the `forward()` method in the MLP class. The MLP has two layers with the following specifications:

- **First layer:** Linear transformation followed by ReLU activation
- **Second layer:** Linear transformation (output layer)

Your implementation should:

- Apply the first linear transformation using `fc1_weight` and `fc1_bias`
- Apply ReLU activation using `self.relu`
- Apply the second linear transformation using `fc2_weight` and `fc2_bias`

Note: Use `torch.mm()` for matrix multiplication.

(2) Implementing Cross Entropy Loss (30 points) Complete the `my_cross_entropy_loss()` function that takes:

- `outputs`: Raw logits from the model `[batch_size, num_classes]`
- `labels`: Ground truth class indices `[batch_size]`

Implementation steps:

- Apply softmax to convert logits to probabilities
- Extract the predicted probability for each correct class
- Calculate negative log likelihood
- Average over the batch

Important:

- Do not modify the random seed.
- Do not import additional libraries.
- Follow the TODOs in the comments of the attached Python file for implementation guidance.

*** END ***